

Greedy Algorithms

- A commonly used paradigm for combinatorial algorithms.
- Informally, in “combinatorial” problems, feasible solutions are subsets of discrete input set, so enumerable in exponential time (say, $O(2^n)$). Greedy algorithms find the optimal by searching only a tiny fraction of this space.
- A precise definition is difficult, but informally an algorithm uses “greedy design principle” if it makes a series of choices, and each choice is locally optimal.
- Why should one expect such a myopic strategy to succeed? Indeed, when greedy strategy works, it says something interesting about the structure(nature) of the problem itself!

Making Change

- The coins in US come in four denominations: 25, 10, 5, 1.
- The “change making” problem is to determine how to convert any amount into minimum number of coins.
- Given an integer $X \in \{0, 1, \dots, 99\}$, find a combination of coins that sum to X using the least number of coins.
- Formally, find integers a, b, c, d with minimum sum $(a+b+c+d)$ so that $X = 25a + 10b + 5c + 1d$

```
In [25]: def makeChange(target: int, coins: list) -> list:
          coins.sort(reverse=True)
          numCoins = []

          for coin in coins:
              numCoins.append({"quantity": target // coin, "coin": coin})
              target -= target // coin * coin

              if not target:
                  break

          if target != 0:
              raise ValueError(
                  "Greedy Algorithm cannot make change with target={} and coins={}"
                  .format(target, coins))

          return numCoins

makeChange(73, [25, 10, 5, 1])
```

```
Out[25]: [{'quantity': 2, 'coin': 25},
           {'quantity': 2, 'coin': 10},
           {'quantity': 0, 'coin': 5},
           {'quantity': 3, 'coin': 1}]
```

Interval Scheduling

- Input: a list of N activities that we want to schedule on a single resource.
- Each activity specified by a start and an end time; only one activity can be scheduled on the resource at a time, and each scheduled activity uses the resource continuously between its start and end time.
- What is the maximum possible number of activities we can schedule?
- Formally, activities is a set $S = \{1, 2, \dots, n\}$, where each activity is specified by its start-end time tuple $(s(i), f(i))$, with $s(i) \leq f(i)$.
- This is a combinatorial problem: output is a subset of $\{1, 2, \dots, n\}$.
- A feasible schedule is a subset in which no two activities overlap.
- Objective: find a feasible schedule of maximum size (number of activities).

Algorithm

- The correct strategy is to process jobs in the Earliest Finish Time order.
- That is, sort the jobs in the increasing order of their finish time. We assume that jobs are given in this order (by simple relabeling): $f(j_1) \leq f(j_2) \leq f(j_3) \dots \leq f(j_n)$

Proof of Correctness

- **Lemma:** For any $i \leq k$, we have that $f(a_i) \leq f(b_i)$. (i.e. i th job in greedy finishes no later than the i th job in the optimal.)
- **Proof:**
 1. True for $i = 1$, by the design of greedy.
 2. Inductively assume this is true for all jobs up to $i - 1$, and prove it for i .
 3. The induction hypothesis says that $f(a_{i-1}) \leq f(b_{i-1})$.
 4. Since $f(b_{i-1}) \leq s(b_i)$, we must also have $f(a_{i-1}) \leq s(b_i)$.
 5. So, the i th job selected by optimal is also available to the greedy as its i th job candidate, so whatever job greedy picks it must have $f(a_i) \leq f(b_i)$.
 6. This proves the lemma.
- **Theorem:** The greedy solution is optimal for the activity selection problem.
- **Proof:**
 1. By contradiction. Suppose A is not optimal, and so OPT must have more jobs than A . That is, $m > k$.
 2. Consider what happens when $i = k$ in our lemma. We have that $f(a_k) \leq f(b_k)$. So, the greedy's last job has finished by the time OPT 's k th job finishes.
 3. If $m > k$, there is some job that optimal accepts after k , and that job is also available to Greedy; it cannot conflict with anything greedy has scheduled.
 4. Because the greedy does not stop until it no longer has any acceptable jobs left, this is a contradiction.

Runtime

- Sorting the jobs takes $O(n \log(n))$.
- After that, the algorithm makes one scan of the list, spending constant time per job = $O(n)$.
- So total time complexity is $O(n \log(n)) + O(n) = O(n \log(n))$.

```
In [37]: def maxActivities(activityList: list) -> dict:
    sortedList = sorted(activityList, key=lambda x: x[1])
    prevEndTime = 0
    activities = list()

    for activity in sortedList:
        if activity[0] >= prevEndTime:
            activities.append(activity)
            prevEndTime = activity[1]

    return {"length" : len(activities), "activities" : activities}

maxActivities([(3,6),(1,4),(4,10),(6,8),(0,2)])
```

```
Out[37]: {'length': 3, 'activities': [(0, 2), (3, 6), (6, 8)]}
```

Interval Partitioning

- Given a set of activities, schedule them all using a minimum number of machines.

Algorithm

- Sort activities by start time.
- Start Room 1 for activity 1.
- For $i = 2$ to n , if activity i can fit in any existing room, schedule it in that room.

Proof of Correctness

- Define depth of input set as the maximum number of activities that are concurrent at any time. Let depth be D .
- Optimal must use at least D rooms because a single room can only house 1 activity and there are D concurrent activities that all need different rooms.
- Greedy uses no more than D rooms because a new room is only created when existing rooms are full, meaning the maximum concurrent amount will be the maximum number of rooms created.

Runtime

- Sorting the jobs takes $O(n \log(n))$.
- After that, the algorithm makes one scan of the list, spending a constant operation to check for an open room, and $O(\log(n))$ operations to insert the a new room, or replace an existing room = $O(n \log(n))$.
- So total time complexity is $O(n \log(n)) + O(n \log(n)) = O(n \log(n))$.

```
In [6]: import heapq

def minPartitions(activityList: list) -> dict:
    if not activityList:
        return 0

    sortedList = sorted(activityList, key=lambda x: x[0])
    endTimes = []
    heapq.heappush(endTimes, sortedList[0][1])

    for i in range(1, len(sortedList)):
        activity = sortedList[i]

        if activity[0] >= endTimes[0]:
            heapq.heappushpop(endTimes, activity[1])
        else:
            heapq.heappush(endTimes, activity[1])

    return {"count": len(endTimes)}

minPartitions([(1,6),(8,13),(15,42),(1,21),(25,31),(35,42)])
```

```
Out[6]: {'count': 2}
```

Huffman Codes

- Goal: encode characters in as few characters as possible
- With variable encoding length, higher frequency characters can be encoded in shorter bitstrings for higher compression
- Prefix Codes: no codeword can be a prefix of another word
- Encode in a binary tree: characters are leaves and branches are bits (path to leaf is binary encoding)
- Huffman codes are only good at encoding static characters. Dynamic data and words have better encoding methods.

Measuring Optimality

- Let C be the input alphabet (set of distinct characters).
- Let $f(p)$ be the frequency of letter p in C .
- Let T be the tree for a prefix code, and $d_T(p)$ the depth of p in T .
- The number of bits (bit complexity) needed to encode our file using this code is:

$$B(T) = \sum_{p \in C} f(p) d_T(p)$$

- We want a code that achieves the minimum possible value of $B(T)$.

Optimal Tree Property: Tree corresponding to optimal code must be full: that is, each internal node has two children. Otherwise we can improve the code.

Huffman's Algorithm

- The algorithm best understood as building the binary tree T that represents its codes.

- Initially, each letter represented by a single-node tree, whose weight equals the letter's frequency.
- Huffman repeatedly chooses the two smallest trees (by weight), and merges them. The new tree's weight is the sum of the two children's weights.
- If there are n letters in the alphabet, there are $n - 1$ merges

Proof of Optimality

- We will use induction on the size of the alphabet $|C|$.
- The base case of $|C| = 2$ is trivial: we have a depth 1 tree, with two leaves, each with code length 1.
- In general, assume induction holds for $|C| = n - 1$, and prove for $|C| = n$.
- Take the last two characters x_{n-1} and x_n , combine them into a single new character z with freq. $f(z) = f(x_{n-1}) + f(x_n)$.
- With x_{n-1}, x_n removed and replaced with z , we have a set of size $|C'| = n - 1$.
- By induction, we find the optimal code tree of C' . This tree has z at some leaf.
- To obtain tree for C , we attach nodes x_{n-1} and x_n as children of z .
- We will show that given optimal tree for C' , this new tree is optimal for C .
- Still one problem: in our construction, the nodes x_{n-1} and x_n will necessarily end up as siblings. (That is, the codes for these two will be identical except in the last bit.)
- How can we choose x_{n-1} and x_n at the onset so that in the optimal tree they are guaranteed to have this property? This is where Huffman's greedy choice enters the proof: we will choose two lowest freq. characters.

Lemma:

- Suppose x and y are two letters of lowest frequency. Then, there exists an optimal prefix code in which codewords for x and y have the same (and maximum) length and they differ only in the last bit.

Proof:

- Start with an optimal prefix code tree T , and modify it so x and y are sibling leaves of max depth, without increasing total cost.
- In modified tree, x and y have the same code length, different only in the last bit.
- Assume optimal tree does not satisfy the claim, and suppose that a and b are the two characters that are sibling leaves of max depth in T .
- Without loss of generality, assume that $f(a) \leq f(b)$ and $f(x) \leq f(y)$
- We have $f(x) \leq f(a)$ and $f(y) \leq f(b)$. (x, y, a, b need not all be distinct.)
- First transform T into T' by swapping the positions of x and a
- Since $d_T(a) \geq d_T(x)$ and $f(a) \geq f(x)$, swap does not increase freq * depth cost:

$$\begin{aligned}
 B(T) - B(T') &= \sum_p [f(p)d_T(p)] - \sum_p [f(p)d_{T'}(p)] \\
 &= [f(x)d_T(x) + f(a)d_T(a)] - [f(x)d_{T'}(x) + f(a)d_{T'}(a)] \\
 &= [f(x)d_T(x) + f(a)d_T(a)] - [f(x)d_T(a) + f(a)d_T(x)] \\
 &= [f(a) - f(x)] * [d_T(a) - d_T(x)] \\
 &\geq 0
 \end{aligned}$$

- Next, transform T' into T'' by exchanging y and b , which also does not increase cost.
- So, we get that $B(T'') \leq B(T') \leq B(T)$. If T was optimal, so is T'' , but in T'' x and y are sibling leaves at the max depth.

Proof of optimality:

- Let T_1 be the optimal tree (induction) for $C + \{z\} - \{x, y\}$.
- We obtain our final tree T by attaching leaves x, y as children of z .
- What is the connection between costs of $B(T)$ and $B(T_1)$?
- For all $p \neq x, y$ depth is the same in both trees, so no difference. For x, y , we have $d_T(x) = d_T(y) = d_{T_1}(z) + 1$. So, the cost increase from modifying T_1 to T is:
 $B(T) - B(T_1) = f(x) + f(y)$ because
 $f(x)d_T(x) + f(y)d_T(y) = [f(x) + f(y)] * [d_{T_1}(z) + 1] = f(z)d_{T_1}(z) + [f(x) + f(y)]$
- The rest of the argument is via contradiction.
- Suppose T is not an optimal prefix code, and another tree T_0 is claimed to be optimal, meaning $B(T_0) < B(T)$.
- By previous lemma, T_0 has x and y as siblings. Imagine replacing parent of x, y with a new leaf z , with freq. $f(z) = f(x) + f(y)$, and call this new tree T'_1 .
- Then, $B(T'_1) = B(T') - f(x) - f(y) < B(T) - f(x) - f(y) < B(T_1)$ which contradicts the claim that T_1 is an optimal prefix code for $C' = C + \{z\} - \{x, y\}$.

Time Complexity

- Time complexity is $O(n \log n)$. Initial sorting plus n heap operations.

In []: `# Insert Code`

Horn Formulas

- Form of boolean logic, and often used in AI systems for logical reasoning.
- Each boolean variable represents an event (or possibility), such as
- x = the murder took place in the kitchen
- y = the butler is innocent
- z = the colonel was asleep at 8pm.
- Recall that Boolean variable can only take one of two values $\{true, false\}$, and a literal is either a variable x or its negation \bar{x}

Constraints among variables represented by two kinds of clauses:

1. Implication: Left-hand-side is an AND of any number of positive literals, and right-hand-side is a single positive literal. $(z \cap u) \rightarrow x$ It asserts that "if the colonel was asleep at 8 pm, and the murder took place at 8pm, then the murder took places in the kitchen." A degenerate statement of the type $\rightarrow x$ means that x is unconditionally true. For instance, "the murder definitely occurred in the kitchen."
 2. Negative: Consists of an OR of any number of negative literals, as in $(\bar{u} \cup \bar{t} \cup \bar{y})$, where $u, t, y, resp.$, means that constable, colonel, and butler is innocent. This clause asserts that "they can't all be innocent."
- A Horn formula is a set of implications and negative clauses.
 - Problem: Given a Horn formula, decide if it is satisfiable, namely, is there an as-ignment of variables so that all clauses are satisfied. Such an assignment is called asatisfying assignment.

Examples:

- The Horn formula $\rightarrow x, \rightarrow y, x \wedge u \rightarrow z, \bar{x} \vee \bar{y} \vee \bar{z}$ has a satisfying assignment $u = 0, x = 1, y = 1, z = 0$.
- But the formula $\rightarrow x, \rightarrow y, x \wedge y \rightarrow z, \bar{x} \vee \bar{y} \vee \bar{z}$ is not satisfiable.

Algorithm

- Brute force approach would take 2^n to account for powerset of inputs.
- The nature of Horn clauses suggests a natural greedy algorithm:
- Initially set all variables to false.
- While there is an unsatisfied Implication clause, set its RHS to true.
- If all pure negative clauses are satisfied, return the assignment; otherwise, formula is not satisfiable.

Correctness Proof

- Clearly, if the algorithm returns a satisfying assignment, then it is a valid assignment because it satisfies all negative and implication clauses.
- To show that if the algorithm does not find a satisfying assignment, there is none, we observe that the algorithm maintains the following invariant. If a certain set of variables is set to true, then they must be true in any satisfying assignment. Namely, we only set a variable true when it is forced upon us.

Time Complexity

- With some care the greedy algorithm can be implemented in linear time (in the length of the formula).

In []: *# Insert Code*

Set Cover

- Input is a (ground) set of n elements $B = \{1, 2, \dots, n\}$ and a collection of m subsets $S = \{S_1, S_2, \dots, S_m\}$, with each $S_i \subseteq B$.
- The problem is to choose the smallest number of subsets whose union is B .
- Example: $B = \{1, 2, 3, 4, 5\}$, and $\{\{1, 2, 3\}, \{2, 4\}, \{3, 4\}, \{4, 5\}\}$. One can cover all items by choosing all four sets, but sets $\{1, 2, 3\}, \{4, 5\}$ suffice.

Algorithm

- Repeat until all elements of B are covered: pick the set S_i containing the largest number of still-uncovered elements.

Runtime

- If the optimal solution uses k sets, the greedy uses $O(k \ln(n))$ sets.

In []: `# Insert Code`

Dijkstra's Algorithm

1. Let S be the set of explored nodes.
2. Let $d(u)$ be the shortest path distance from s to u , for each $u \in S$.
3. Initially $S = \{s\}$, $d(s) = 0$, and $d(u) = \infty$, for all $u \neq s$.
4. While $S \neq V$ do
5. Select $v \notin S$ with the minimum value of $d'(v) = \min_{(u,v), u \in S} d(u) + \text{cost}(u, v)$
6. Add v to S , set $d(v) = d'(v)$.

Correctness Proof

1. Argue that at any time $d(v)$ is the shortest path distance to v , for all $v \in S$.
2. Consider the instant when node v is chosen by the algorithm. Let (u, v) be the edge, with $u \in S$, that is incident to v .
3. Suppose, for the sake of contradiction, that $d(u) + \text{cost}(u, v)$ is not the shortest path distance to v . Instead a shorter path P exists to v .
4. Since that path starts at s , it has to leave S at some node. Let x be that node, and let $y \notin S$ be the edge that goes from S to \bar{S} .
5. So our claim is that $\text{length}(P) = d(x) + \text{cost}(x, y) + \text{length}(y, v)$ is shorter than $d(u) + \text{cost}(u, v)$. But note that the algorithm chose v over y , so it must be that $d(u) + \text{cost}(u, v) \leq d(x) + \text{cost}(x, y)$.
6. In addition, since $\text{length}(y, v) > 0$, this contradicts our hypothesis that P is shorter than $d(u) + \text{cost}(u, v)$.
7. Thus, the $d(v) = d(u) + \text{cost}(u, v)$ is correct shortest path distance.

In []: `# Insert Code`

Kruskal's Algorithm

1. If the shortest edge connects two previously unconnected vertices, add that edge to the spanning tree.
2. Continue repeating step 1 until all the vertices are connected.

Correctness Proof

1. For simplicity, assume that all edge costs are distinct so that the MST is unique. Otherwise, add a tie-breaking rule to consistently order the edges.
2. Proof by contradiction: let (v, w) be the first edge chosen by Kruskal that is not in the optimal MST.
3. Consider the state of the Kruskal just before (v, w) is considered.
4. Let S be the set of nodes connected to v by a path in this graph. Clearly, $w \notin S$.

5. The optimal MST does not contain (v, w) but must contain a path connecting v to w , by virtue of being spanning.
6. Since $v \in S$ and $w \notin S$, this path must contain at least one edge (x, y) with $x \in S$ and $y \notin S$.
7. Note that (x, y) cannot be in Kruskal's graph at the time (v, w) was considered because otherwise y will have been in S .
8. Thus, (x, y) is more expensive than (v, w) because it came after (v, w) in Kruskal's scan order.
9. If we replace (x, y) with (v, w) in the optimal MST, it remains spanning and has lower cost, which contradicts its optimality.
10. So, the hypothesis that (v, w) is not in optimal must be false.

In []: *# Insert Code*